

AADHI DHARMAR

Agentic AI Engineer • GenAI & RAG Systems • AI Security

Coimbatore, India | +91 6380072080 | dharmaraadhi@gmail.com | github.com/Aadhidharmar001 |
linkedin.com/in/aadhi-dharmar-a35679293

PROFESSIONAL SUMMARY

Final-year B.Tech AI & Data Science student building production-quality agentic AI workflows, action guardrail gateways, and stateful RAG systems. Proficient in LangGraph orchestration, FastAPI backend microservices, vector databases, and containerized AWS cloud infrastructure (ECS Fargate, PostgreSQL, Docker). Built production-ready reverse proxies and YAML security policy engines to govern autonomous AI agent tool execution.

TECHNICAL SKILLS

Languages & Backend:	Python, SQL, FastAPI, PostgreSQL, SQLAlchemy, REST APIs, Server-Sent Events (SSE)
GenAI & Agentic Systems:	LangGraph, LangChain, RAG Systems, ChromaDB, OpenAI API, Groq API, PyTorch, HuggingFace, NLP
Cloud, DevOps & Tools:	AWS (ECS Fargate, CloudWatch), Docker, pytest, Git/GitHub, Streamlit
Security & Governance:	AI Agent Security & WAF, Action Guardrails, YAML Security Policies, Rate Limiting, SQL Injection Detection, HITL

KEY PROJECTS

GuardWAF — Enterprise AI Agent Guardrail Gateway (2026)

FastAPI • AWS ECS Fargate • Docker • PostgreSQL • SQLAlchemy • CloudWatch • Grok/OpenAI API

- Architected a production-ready AI Agent Web Application Firewall (WAF) reverse proxy intercepting, governing, and auditing tool/API calls before execution to prevent unauthorized or out-of-scope LLM agent actions.
- Engineered a sliding-window rate-limiting algorithm to block runaway agent loops and stateful sequence validation middleware to enforce secure multi-step execution workflows.
- Implemented session-aware data scope verification, parameter validation, and SQL injection detection filters using a configurable YAML-driven security policy engine.
- Designed a Human-in-the-Loop (HITL) approval workflow featuring automated risk scoring for sensitive operations, paired with real-time audit logging and an SSE monitoring dashboard.
- Containerized microservices using Docker and deployed to AWS ECS Fargate with PostgreSQL audit trail persistence and Amazon CloudWatch operational logging.

NeuralRAG — Agentic Retrieval-Augmented Generation System (2025 – 2026)

LangGraph • ChromaDB • OpenAI / Groq API • HuggingFace • Streamlit • Python

- Architected an agentic RAG system using LangGraph stateful graphs for multi-step reasoning, document chunking, and ChromaDB vector retrieval to ground LLM responses in verifiable source documents.
- Developed a domain-filtered web search engine with JSON state persistence and source-transparency provenance labels, enabling complete auditability for every generated answer.
- Engineered a query-routing agent for summary-intent classification and refactored the Streamlit frontend into a modular, production-ready architecture, optimizing response latency through efficient retrieval.

Intelligent Search Engine (2025)

Python • FastAPI • Vector Embeddings • Semantic Search • Render

- Developed a semantic search engine leveraging vector embeddings and similarity ranking to surface contextually relevant results beyond exact keyword matching.
- Engineered a lightweight FastAPI REST microservice and deployed the production backend to Render, improving responsiveness using vector search and optimized endpoints.

EDUCATION

B.Tech in Artificial Intelligence & Data Science — Sri Shakthi Institute of Engineering and Technology (Expected May 2027)
Relevant Coursework: Data Structures & Algorithms, Machine Learning, Deep Learning, DBMS, NLP, Software Engineering

CERTIFICATIONS & ACHIEVEMENTS

- **Prompt Engineering** — Simplilearn SkillUp (Feb 2026) | **Intro to Data Science & NLP** — Infosys Springboard (Jun 2025)
 - **Data Visualization** — Tata Forage (Aug 2025) | **Career Essentials in Generative AI** — Microsoft & LinkedIn (Aug 2024)
 - **AI for All (AI Aware & Appreciate)** — Intel & Govt. of India (Jul 2025)
- **Secured 15th Place Nationally at VIT Hackathon (AI/ML Track) out of 100+ competing teams by engineering an AI Medical Diagnostic Chatbot achieving 86% accuracy within 24 hours.**